

Extra Credit - Deployment

Assume you are asked to get your code running in the cloud using AWS. What tools and AWS services would you use to deploy the API, database, and a scheduled version of your data ingestion code? Write up a description of your approach.

1. Database Deployment

Service Used: Amazon RDS (Relational Database Service)

Steps:

- **Create RDS Instance:**
 - Use Amazon RDS to deploy a PostgreSQL database instance. Choose the appropriate instance type based on your workload requirements (e.g., t3.medium for moderate usage).
 - Configure the database instance details such as instance size, storage type (SSD-backed), allocated storage size, and availability zone.
 - Set up the database engine version (PostgreSQL) and optionally enable automatic minor version upgrades.
 - Configure security settings:
 - **Security Group:** Create or select a security group that allows inbound traffic on port 5432 (PostgreSQL default port) from your Elastic Beanstalk environment and Lambda function.
 - **Subnet Group:** Place the database instance in a custom subnet group within your Virtual Private Cloud (VPC) for network isolation.
- **Integrate with AWS Secrets Manager:**
 - Store database credentials securely in AWS Secrets Manager. Create a secret named `weather_db_credentials` with key-value pairs for `username` and `password`.
 - Ensure IAM policies are configured to grant necessary permissions for accessing the secret.
- **Configure Application Environment:**
 - Modify your application's configuration to fetch database credentials securely from AWS Secrets Manager at runtime.
 - Update `SQLALCHEMY_DATABASE_URI` in your application's environment variables to reference the secret stored in AWS Secrets Manager.

2. API Deployment

Service Used: AWS Elastic Beanstalk

Steps:

- **Create Elastic Beanstalk Application:**
 - Set up an Elastic Beanstalk environment for deploying and managing your Flask-based API application.
 - Choose the appropriate platform (Python) and environment type (Web server environment).
 - Configure environment properties:
 - **Environment Variables:** Define environment variables such as `SQLALCHEMY_DATABASE_URI` pointing to the RDS instance endpoint, database name, username, and password.
 - **Instance Configuration:** Configure instance types, scaling options (e.g., load balancing, auto-scaling), and EC2 instance security settings.
 - **Network Configuration:** Associate the environment with your VPC and subnet configuration. Ensure security groups allow inbound HTTP traffic (port 80) from the internet.
- **Deploy Flask Application:**
 - Package your Flask application into a ZIP file containing `app.py`, `requirements.txt` listing dependencies (including `Flask`, `SQLAlchemy`, `psycopg2`), and other necessary files.
 - Upload the ZIP file to Elastic Beanstalk using the AWS Management Console or AWS CLI.
 - Monitor deployment status through Elastic Beanstalk console and view logs for troubleshooting (`eb logs` command).
- **Access AWS Secrets Manager in Code:**
 - Use AWS SDKs (`boto3` for Python) within your Flask application (`app.py`) to fetch database credentials securely from AWS Secrets Manager.
 - Retrieve the database credentials (`username` and `password`) and construct `SQLALCHEMY_DATABASE_URI` dynamically based on the fetched values.
- **Configure Auto-scaling and Load Balancing:**
 - Elastic Beanstalk automatically provisions an Application Load Balancer (ALB) to distribute incoming traffic across multiple EC2 instances running your Flask application.
 - Configure auto-scaling policies based on CPU utilization or request metrics to automatically add or remove EC2 instances to handle varying traffic loads.

3. Scheduled Data Ingestion

Service Used: AWS Lambda with EventBridge (formerly CloudWatch Events)

Steps:

- **Create AWS Lambda Function:**
 - Develop a Python function (`data_ingestion_lambda.py`) using the `boto3` library to interact with AWS services and process data ingestion tasks.
 - Ensure the Lambda function has appropriate permissions:
 - **Execution Role:** Create an IAM role granting permissions to access AWS services like RDS (for database operations) and CloudWatch Logs (for logging).
 - **Environment Variables:** Set environment variables such as `SQLALCHEMY_DATABASE_URI` to securely connect to the RDS instance.
- **Use AWS Secrets Manager in Lambda Function:**
 - Grant Lambda function permissions (`secretsmanager:GetSecretValue`) to access the `weather_db_credentials` secret stored in AWS Secrets Manager.
 - Update the Lambda function code (`data_ingestion_lambda.py`) to retrieve database credentials dynamically from AWS Secrets Manager.
- **Upload Code to Lambda:**
 - Package your Lambda function code (`data_ingestion_lambda.py`) along with any additional dependencies (`psycopg2`, `sqlalchemy`) into a ZIP file.
 - Upload the ZIP file to Lambda using the AWS Management Console, AWS CLI, or AWS SDKs.
- **Schedule Execution using EventBridge:**
 - Use AWS EventBridge to schedule the Lambda function to execute at predefined intervals (e.g., daily at 2 AM UTC).
 - Configure the schedule using cron expressions (`cron(0 2 * * ? *)`) to trigger the Lambda function based on UTC time.
- **Error Handling and Logging:**
 - Implement error handling within the Lambda function to capture exceptions and errors encountered during data ingestion.
 - Log detailed execution logs to AWS CloudWatch Logs for monitoring and troubleshooting (`print` statements or `logging` module).

4. Additional Considerations

- **Security:**
 - **IAM Roles and Policies:** Apply the principle of least privilege by creating IAM roles with specific permissions for each service (RDS, Lambda, Elastic Beanstalk).
 - **Encryption:** Enable encryption at rest for the RDS instance using AWS KMS (Key Management Service) and in-transit using SSL/TLS.
 - **Network Isolation:** Use VPC Endpoints for accessing AWS services (e.g., S3 for data storage) privately without exposing them to the internet.
- **Monitoring and Logging:**
 - **CloudWatch Metrics and Alarms:** Set up CloudWatch alarms to monitor API latency, error rates, and Lambda function invocations. Create custom metrics for monitoring specific application-level metrics.
 - **Dashboards:** Build custom CloudWatch dashboards to visualize key performance indicators (KPIs) and operational metrics for the API and data ingestion processes.
- **Scaling:**
 - **Elastic Beanstalk Auto-scaling:** Configure Elastic Beanstalk environment to scale automatically based on CPU utilization or request metrics.
 - **Lambda Concurrency:** Adjust Lambda function concurrency limits to handle concurrent executions during peak data ingestion periods.
- **Backup and Recovery:**
 - **RDS Automated Backups:** Enable automated backups for the RDS instance to retain backups for point-in-time recovery.
 - **Snapshot Management:** Create manual snapshots of the RDS instance before making significant changes to the database schema or configuration.